

# CS189/289A – Spring 2017 — Homework 5

Yicheng Chen, SID 26943685

Collaborators: NONE)

## 1. Implement decision trees

Code:

```
class decision_tree:
```

```
    class node:
```

```
        def __init__(self, left, right, split_rule, is_leaf, label):
            self.left = left
            self.right = right
            self.split_rule = split_rule
            self.is_leaf = is_leaf
            self.label = label
```

```
    def __init__(self, max_depth=1e10):
        self.max_depth = max_depth
```

```
    def max_count(self, array):
        return max(set(array), key=array.tolist().count)
```

```
    # utility function for entropy calculation
```

```
    def entropy(self, indices):
        p = itemfreq(a[:, 1]).astype(float) / len(indices)
        return -p.dot(np.log2(p))
```

```
    # calculate entropy the number of instances in each class in known
```

```
    def entropy_n(self, all_n):
        p = all_n / (np.sum(all_n)+1e-20)
        return -p.dot(np.log2(p+1e-20))
```

```
    # calculate the impurity("badness") of the specified split on the input data
```

```
    def impurity(self, left_label_hist, right_label_hist):
        Sl = np.sum(left_label_hist)
        Sr = np.sum(right_label_hist)
        return (Sl*self.entropy_n(left_label_hist) + Sr * self.entropy_n(right_label_hist))
```

```
    # find the threshold that best split the data points with a certain feature
```

```

# Note: <= th goes to S_left and > th goes to S_right
def find_threshold(self, feature, labels):
    all_f = sorted(set(feature)) # sorted in ascending order
    all_l = set(labels) # list unique labels

    freq_mat = np.zeros([len(all_f), len(all_l)])
    for i, f in enumerate(all_f):
        for j, l in enumerate(all_l):
            freq_mat[i, j] = len(labels[np.where(labels[np.where(feature==f)]==l)])

    # calculate the average of two neighboring values as threshold
    # iterates from min to max
    all_threshold = (np.hstack((all_f[1:], all_f[-1]))) + all_f) / 2.

    # in the beginning, all goes to the right node
    n_left = np.zeros([len(all_l)])
    n_right = np.sum(freq_mat, axis=0)
    n_left_sum = 0
    min_threshold = all_threshold[0]
    min_H = self.entropy(n_left, n_right)
    # loop through all threshold to find the one with the minimum entropy
    for i, th in enumerate(all_threshold):
        n_left += freq_mat[i, :]
        n_right -= freq_mat[i, :]
        H = self.entropy(n_left, n_right)
        if H < min_H:
            min_H = H
            min_threshold = th
    return min_threshold, min_H

# find the best feature and threshold to split data points
def segmenter(self, data, labels, m=-1):
    d = data.shape[1]
    if m == -1:
        all_features = np.arange(d)
    else:
        all_features = np.random.choice(range(d), m, replace=False)
    min_H = 1e20
    min_th = 0
    min_i = 0
    for i in all_features:
        threshold, H = self.find_threshold(data[:, i], labels)
        if H < min_H:
            min_H = H
            min_th = threshold
            min_i = i
    return min_i, min_th

```

```

# the recurrence function that builds the decision tree
def grow_tree(self, S, depth, m=-1):
    if len(set(self.labels[S])) == 1 or depth >= self.max_depth: # pure node or reach m
        return self.node(left=None, right=None, split_rule=None, is_leaf=1, \
                        label=self.max_count(self.labels[S]))
    else:
        min_i, min_th = self.segmenter(self.data[S, :], self.labels[S], m=m)
        # the following comprehension might be slow
        # S1 = [j for j, x in enumerate(self.data[:, min_i]) if x <= min_th and j in S]
        # Sr = [j for j, x in enumerate(self.data[:, min_i]) if x > min_th and j in S]
        # update: faster:
        S1 = [j for j in S if self.data[j, min_i] <= min_th]
        Sr = [j for j in S if self.data[j, min_i] > min_th]
        # another: faster:
        # S1 = np.intersect1d(np.where(self.data[:, min_i] <= min_th), S)
        # Sr = np.intersect1d(np.where(self.data[:, min_i] > min_th), S)
        if len(S1) == 0 or len(Sr) == 0:
            return self.node(left=None, right=None, split_rule=None, is_leaf=1, \
                            label=self.max_count(self.labels[S]))
        else:
            return self.node(left=self.grow_tree(S1, depth+1, m), right=self.grow_tree(
                Sr, depth+1, m), split_rule = (min_i, min_th), \
                            is_leaf=0, label=None)

# train the decision tree
def train(self, data, labels, m=-1):
    self.data = data
    self.labels = labels
    S = np.array(range(len(labels)))
    self.root = self.grow_tree(S, 1, m=m)
    self.data = None
    self.labels = None
    return self

# predict labels of test data
def predict(self, data, verbose=False):
    if data.ndim == 1: # special case of only 1 row (it becomes a 1d vector in numpy)
        data = np.reshape(data, [1, len(data)])
        N = 1
    else:
        N = data.shape[0]
    labels = np.zeros(N)

    # predict each data point
    for i in range(N):
        d = data[i, :]
        current_node = self.root

```

```
# going down along the tree
while not current_node.is_leaf: # not reach leaf yet
    idx = current_node.split_rule[0]
    th = current_node.split_rule[1]
    if d[idx] <= th:
        current_node = current_node.left
        if verbose:
            print('Going left')
    else:
        current_node = current_node.right
        if verbose:
            print('Going right')
if verbose:
    print()

labels[i] = current_node.label
return labels

# calculate the prediction accuracy
def accuracy(self, data, true_labels):
    labels = self.predict(data, verbose=False)
    N = len(labels)
    return np.sum(labels == true_labels) / float(N)
```

## 2. Implement random forests

Code:

```
class random_forest:
    def __init__(self, n_trees=20, n_sample=1000, n_feature=-1, max_depth=1e10):
        self.n_trees = n_trees
        self.n_sample = n_sample
        self.n_feature = -1
        self.max_depth = max_depth
        self.trees = np.array([decision_tree(max_depth)] * n_trees)

    def train_parallel(self, data, labels, verbose=True):
        if self.n_feature == -1:
            d = data.shape[1]
            self.n_feature = int(np.sqrt(d)) # num of random features = sqrt(d) is a good g
        pool = mp.Pool()
        results = np.zeros(self.n_trees, dtype=object)
        print('# Trees = %d' % self.n_trees)
        for i, dt in enumerate(self.trees):
            #
            print(i)
            idx = np.random.choice(range(len(data)), self.n_sample)
            sub_data = data[idx, :]
            sub_labels = labels[idx]
            results[i] = pool.apply_async(dt.train, args=(sub_data, sub_labels, self.n_feature))
            if verbose:
                print('%d: Tree loaded.' % i)
        #
        dt.train(sub_data, sub_labels, m=self.n_feature) # activate the random feature
        self.trees = [p.get() for p in results]
        print('Trees trained.')
        pool.close()

    def train(self, data, labels):
        if self.n_feature == -1:
            d = data.shape[1]
            self.n_feature = int(np.sqrt(d)) # num of random features = sqrt(d) is a good g
        print('# Trees = %d' % self.n_trees)
        for i, dt in enumerate(self.trees):
            print(i)
            idx = np.random.choice(range(len(data)), self.n_sample)
            sub_data = data[idx, :]
            sub_labels = labels[idx]
            dt.train(sub_data, sub_labels, m=self.n_feature) # activate the random feature

    def predict(self, data, verbose=False):
        if data.ndim == 1: # special case of only 1 row (it becomes a 1d vector in numpy)
            data = np.reshape(data, [1, len(data)])
            N = 1
```

```
    else:
        N = data.shape[0]
        labels = np.zeros(N)

    # predict each data point
    for i in range(N):
        votes = np.zeros(self.n_trees)
        for j, t in enumerate(self.trees):
            votes[j] = t.predict(data[i, :], verbose=verbose)
        labels[i] = t.max_count(votes)
    return labels

def accuracy(self, data, true_labels):
    labels = self.predict(data, verbose=False)
    N = len(labels)
    return np.sum(labels == true_labels) / float(N)
```

### 3. Describe implementation details

- (a) Categorical features are vectorized. Missing values are replaced by the most frequent element or mean value.
- (b) The stopping criteria is the maximum depth of tree.
- (c) Nothing special to speed up decision tree. But for random forest, the multiprocessing module is used to train trees in parallel.
- (d) The random forest is basically a collection of decision trees with random data sampling and feature sampling. Only a wrapper class and small modification of the decision tree class are needed.
- (e) I think the parallel part is quite cool!

## 4. Performance evaluation

(a) Decision tree:

Data	Training Accuracies	Validation Accuracies
Spam	0.8071	0.7848
Census	0.8924	0.8400
Titanic	0.9237	0.77

Random forest:

Data	Training Accuracies	Validation Accuracies
Spam	0.7979	0.7816
Census	0.8924	0.8592
Titanic	0.9237	0.87

(b) Kaggle:

Display name: Yicheng Chen

Public scores:

Data	Public score
Spam	0.97880
Census	0.86450
Titanic	0.81290



## 5. Writeup for the spam dataset

- (a) Bag-of-words and TF-IDF features are used for the spam dataset.
- (b) The first data point of the training set (max depth=15):

```

Feature 28 <= 0.500000
Feature 29 <= 0.500000
Feature 19 <= 0.500000
Feature 25 <= 1.500000
Feature 7 <= 0.500000
Feature 13 <= 0.500000
Feature 3 <= 0.500000
Feature 0 <= 0.500000
Feature 26 <= 0.500000
Feature 31 <= 0.500000
Feature 6 <= 0.500000
Feature 20 <= 0.500000
Feature 16 <= 0.500000
Feature 30 <= 0.500000
Classification: 0

```

- (c) Random forest split at root (20 trees):

```

Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 16 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000
Feature 19 <= 0.500000

```

Feature 19 <= 0.5 (19 trees)  
 Feature 16 <= 0.5 (1 tree)

## 6. Writeup for the census dataset


- (a) The 2nd order polynomial features are used.
- (b) The first data point of the training set (max depth=15):

```
Feature 54 <= 0.500000
Feature 4 <= 7055.500000
Feature 21 <= 488.000000
Feature 12 > 1332.500000
Feature 5 <= 2218.500000
Feature 84 > 0.500000
Feature 9 <= 394.000000
Feature 1 > 32.500000
Feature 8 > 11725338.500000
Classification: 0
```

- (c) Random forest split at root (20 trees):

```
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
Feature 48 <= 0.500000
```

- (d) The accuracy starts low because the bias is high when the trees are limited in depth. But it then rises because the bias is reduced with deeper trees.



## 7. Writeup for the Titanic dataset

The tree looks like this:

