

CS189/289A – Spring 2017 — Homework 6

Yicheng Chen, SID 26943685

1. Derivations

$$\begin{aligned}\nabla_{W_i} L &= \frac{\partial L}{\partial z_i} \frac{\partial z_i}{\partial W_i} \\ &= \frac{\partial L}{\partial z_i} s'(W_i \cdot h) h\end{aligned}$$


Where $s'(W_i \cdot h) = \nabla_{W_i} \sigma(W_i \cdot h)$ and h is the hidden unit output.

$$\begin{aligned}\nabla_{V_i} L &= \frac{\partial L}{\partial h_i} \frac{\partial h_i}{\partial V_i} \\ &= \sum_{j=1}^k \frac{\partial L}{\partial z_j} \frac{\partial z_j}{\partial h_i} t'(V_i \cdot x) x\end{aligned}$$

Where $t'(V_i \cdot x) = \nabla_{V_i} \tanh(V_i \cdot x)$ and x is a random training sample.

2. Implementation


- (a) Hyperparameters tuned: the step size and the number of hidden layer neurons.
- (b) Training Accuracy: 0.9508
- (c) Validation Accuracy: 0.8578
- (d) Plot of the loss value versus number of epochs:




- (e) For Kaggle submission, HOG feature is used for model training and predicting. Kaggle Score: 0.91163
Display Name: Yicheng Chen
- (f) Code: See appendix

3. Visualization

(a) 5 digits in the validation set that are correctly classified:



(b) 5 digits in the validation set that are not correctly classified:



4. Bells and Whistles

Due to the time constraint, no bells and whistles were implemented.

Appendix: Code

```

class neural_network():

    def __init__(self, layer_sizes=[784, 200, 26], step_size=0.1):
        self.layer_sizes = layer_sizes
        self.step_size=step_size #??
        # support classical 3-layer network only
        self.V = np.random.normal(loc=0, scale=1/np.sqrt(layer_sizes[0]+1), size=(layer_sizes[0], layer_sizes[1]))
        self.W = np.random.normal(loc=0, scale=1/np.sqrt(layer_sizes[1]+1), size=(layer_sizes[1], layer_sizes[2]))

    # activation function: tanh
    def tanh(self, x):
        x = np.minimum(500, x)
        x = np.maximum(-500, x)
        a = np.exp(x)
        b = np.exp(-x)
        return (a-b)/(a+b)

    # derivative of tanh function
    def tanh_d(self, x):
        t = self.tanh(x)
        return 1 - np.multiply(t, t)

    # activation function: sigmoid
    def sigmoid(self, x):
        return 1/(1+np.exp(-x))

    # derivative of sigmoid function
    def sigmoid_d(self, x):
        s = self.sigmoid(x)
        return np.multiply(s, 1-s)

    # activation function: ReLU
    def relu(self, x):
        return np.maximum(0, x)

    # derivative of ReLU function
    def relu_d(self, x):
        return (x>0).astype(float)

    # activation function: softmax
    def softmax(self, x):
        e = np.exp(x)
        return e / np.sum(e)

    # loss function: cross-entropy
    def cross_entropy(self, z, y):

```

```

        return -np.sum(np.multiply(y, np.log(z+1e-20)) + np.multiply((1-y), np.log(1-z+1e-20))

# derivative of cross-entropy loss function
def cross_entropy_d(self, z, y):
    return -np.divide(y, z+1e-20) + np.divide((1-y), (1-z+1e-20))

# calculate h = s(Vx) and z = s(Wh), aka forward propagation
def forward_prop(self, x):
    x1 = np.append(x, 1) # 1 at the end
    h = self.sigmoid(self.V.dot(x1)) # tanh for hidden layer
    h1 = np.append(h, 1) # 1 at the end
    z = self.sigmoid(self.W.dot(h1)) # sigmoid for output layer
    return x1, h1, z

# calculate the gradient of weight matrices, aka backward propagation
def backward_prop(self, x1, h1, z, y):
    dLdz = self.cross_entropy_d(z, y)
    s2d = self.sigmoid_d(self.W.dot(h1))
    # the gradient of W
    dW = np.outer(np.multiply(dLdz, s2d), h1)

    s1d = self.sigmoid_d(self.V.dot(x1))
    dLdh1 = np.sum(np.dot(np.diag(np.multiply(dLdz, s2d)), self.W), axis=0)
    dLdh = dLdh1[:-1] # don't include the 1
    # the gradient of V
    dV = np.outer(np.multiply(dLdh, s1d), x1)
    return dV, dW

# vectorize data label
def vectorize_y(self, y):
    M = np.max(y)
    N = len(y)
    yv = np.zeros([N, M])
    for i in range(N):
        yv[i, y[i]-1] = 1
    return yv

# transform vectorized y back to labels
def labelize_y(self, y):
    return np.argmax(y, axis=1) + 1

# train NN using stochastic gradient descent
# X: Nxd data matrix
# y: Nx1 data label, range 1~26
def train_sgd(self, X, y, N_epoch=10):
    all_loss = []
    yv = self.vectorize_y(y)
    for epoch in range(0, N_epoch):

```

```

        print('epoch #%d' % epoch)
        for i in range(X.shape[0]):
            xr = X[i, :]
            yr = yv[i, :]
            # FP
            x1, h1, z = self.forward_prop(xr)
            # BP
            dV, dW = self.backward_prop(x1, h1, z, yr)
            # GD
            self.V -= self.step_size * dV
            self.W -= self.step_size * dW
        print(' Training Accuracy: %f' % self.accuracy(X, y))
        all_loss.append(self.loss(X, yv))
    return all_loss

def predict(self, X):
    z = np.zeros([len(X), len(self.W)])
    for i, x in enumerate(X):
        _, _, z[i, :] = self.forward_prop(x)
    z = np.argmax(z, axis=1) + 1
    return z

def accuracy(self, X, y):
    z = self.predict(X)
    N = len(y)
    return np.sum(z == y) / float(N)

def loss(self, X, y):
    loss = 0
    z = np.zeros([len(X), len(self.W)])
    for i, x in enumerate(X):
        _, _, z[i, :] = self.forward_prop(x)
    for i in range(len(z)):
        yi = y[i, :]
        zi = z[i, :]
        loss += self.cross_entropy(zi, yi)
    return loss / len(z)

```